


Deep Learning

1. What is Deep Learning?

- **Subset of Machine Learning:** Deep Learning (DL) is a specialized subfield of Machine Learning (ML), which itself is a subfield of Artificial Intelligence (AI).
- **Inspired by the Human Brain:** It uses artificial neural networks with multiple layers (hence "deep") to learn and make decisions, drawing inspiration from the structure and function of the human brain's interconnected neurons.
- **Learning Representations:** Instead of being explicitly programmed, deep learning models learn hierarchical representations of data. Lower layers learn simple features, and higher layers combine these to learn more complex concepts.



Key Components

1. Neurons and Layers

- **Neuron:** Basic unit that receives inputs, applies a weight and bias, passes through an activation function.
- **Layers:**
 - **Input Layer:** Receives raw features
 - **Hidden Layers:** Process and transform inputs
 - **Output Layer:** Produces final predictions

2. Activation Functions

Introduce non-linearity. Common ones:

- ReLU (Rectified Linear Unit)
- Sigmoid
- Tanh

- Softmax (for classification)

3. Loss Function

Measures the error between predicted and true values.

- Examples: MSE (regression), Cross-Entropy (classification)

4. Optimization

Backpropagation + Gradient Descent to adjust weights

- Optimizers: SGD, Adam, RMSprop

Popular Architectures

| Type | Description | Example Use |
|------------------------|-----------------------------------|----------------------|
| Feedforward NN (FNN) | Basic dense layers | Tabular data |
| Convolutional NN (CNN) | Spatial hierarchies | Image classification |
| Recurrent NN (RNN) | Sequence modeling | Language modeling |
| Transformers | Attention-based, parallel | NLP, Vision |
| Autoencoders | Encoding/decoding for compression | Anomaly detection |
| GANs | Generator vs Discriminator | Image synthesis |

Frameworks & Tools

| Framework | Language | Strength |
|------------|----------|------------------------------|
| TensorFlow | Python | Industry-grade, scalable |
| PyTorch | Python | Research-friendly, dynamic |
| Keras | Python | High-level, easy prototyping |
| ONNX | Interop | Model deployment/interchange |



Use Cases

- **Computer Vision:** Image classification, detection, segmentation
- **NLP:** Sentiment analysis, translation, chatbots
- **Speech:** Recognition, synthesis
- **Recommendation Systems**
- **Medical Diagnosis, Finance, Autonomous Vehicles**

Core Architectural Paradigms (High-Level):

- **Artificial Neural Networks (ANNs):**
 - **Concept:** The foundational building blocks. Consist of interconnected nodes (neurons) organized in layers (input, hidden, output). Each connection has a weight that is learned during training.
 - **What we study:** Neuron models (e.g., perceptron), activation functions (Sigmoid, ReLU, Tanh), feedforward propagation, backpropagation (for learning weights), loss functions, and optimization algorithms (e.g., Gradient Descent and its variants).
 - **Relevance:** Basis for all other deep learning architectures. Used for simpler classification and regression tasks on structured data.
- **Convolutional Neural Networks (CNNs or ConvNets):**
 - **Concept:** Specialized for processing grid-like data, such as images. Employ operations like convolution (to detect local features like edges, textures) and pooling (to downsample and create invariance).

- **What we study:** Convolutional layers, filters/kernels, pooling layers (max, average), padding, stride, receptive fields, and well-known architectures (e.g., LeNet, AlexNet, VGG, ResNet, Inception).
- **Relevance:** Dominant in computer vision tasks (image classification, object detection, segmentation), but also applied to other domains like NLP.
- **Recurrent Neural Networks (RNNs):**
 - **Concept:** Designed to handle sequential data (e.g., text, time series, speech) by having connections that form directed cycles, allowing information from previous steps to persist (memory).
 - **What we study:** Handling sequences, vanishing/exploding gradients problem, Long Short-Term Memory (LSTM) units, Gated Recurrent Units (GRU), sequence-to-sequence models, attention mechanisms (often studied alongside RNNs/Transformers).
 - **Relevance:** Crucial for Natural Language Processing (machine translation, sentiment analysis), speech recognition, time series forecasting.
- **Transformers:**
 - **Concept:** A more recent architecture (emerged prominently around 2017) that has revolutionized NLP and is increasingly used in other domains like vision. Relies heavily on "self-attention" mechanisms to weigh the importance of different parts of the input sequence, allowing for parallel processing of sequences (unlike traditional RNNs).
 - **What we study:** Self-attention, multi-head attention, positional encodings, encoder-decoder structures (e.g., BERT, GPT).
 - **Relevance:** State-of-the-art for many NLP tasks and showing great promise in computer vision (Vision Transformers - ViT) and other areas.

5. Key Enabling Concepts & Tools:

- **Transfer Learning:** Reusing a pre-trained model (trained on a large benchmark dataset like ImageNet or a massive text corpus) as a starting point for a new, related task. This is highly effective when you have limited data for your specific problem.
- **Embeddings:** Dense vector representations of discrete inputs (like words or categories) learned by the network, capturing semantic relationships.
- **Regularization Techniques:** Methods (e.g., L1/L2 regularization, Dropout, Batch Normalization) to prevent overfitting, where the model performs well on training data but poorly on unseen data.
- **Optimization Algorithms:** Advanced optimizers beyond basic SGD, such as Adam, RMSprop, AdaGrad, which help in efficiently navigating the complex loss landscapes of deep networks.
- **Hyperparameter Tuning:** The process of finding the optimal set of hyperparameters (e.g., learning rate, number of layers, number of neurons per layer) for a model, often requiring systematic experimentation.

Deep Learning

- ① ANN → Artificial Neural N/w $\left\{ \begin{array}{l} \rightarrow \text{Classification} \\ \rightarrow \text{Regression} \end{array} \right\}$
- ② CNN → Convolutional Neural N/w → I/P: Images, Video frames → RNN, MASKED RNN, Detection, YOLO V5, V6, V7
- ③ RNN → Recurrent Neural N/w → NLP → NLP, Time Series
I/P: Text, Time Series

FRAMEWORK

TENSORFLOW

End to End Project

{ Word Embedding, LSTM RNN, GRU RNN,
Bidirectional LSTM RNN, Encoder Decoder,
Transformers, BERT }

② Why Deep Learning Is Becoming popular?

2005 → Facebook, Youtube, Whatsapp, Linked, Twitter } ⇒ Social Media platform

⋮

2011-2012

DATA WAS GENERATED

Exponentially

Images, Text, Documents,

Videos

→ **Big DATA** → HADOOP → Unstructured DATA

↓

2011 ⇒ Cloudera, Hooten works

2011-2012 :

① Hardware Requirement → CPU's cost ↓ ↓ → NVIDIA GPU
TRAIN MODELS Price ↓ ↓


② Huge amount of data is getting generated → Deep Learning
Models Perform well

Netflix → movie streaming

↓ platform

Recommendation system

↓
Increase Revenue




The Perceptron

1. Introduction & History:

- The **Perceptron** is the simplest form of a neural network, a foundational algorithm for **supervised learning** of **binary classifiers**.
- It's a type of **linear classifier**, meaning it learns a linear decision boundary to separate two classes.
- It's inspired by a single biological neuron, aiming to mimic basic information processing.

2. Analogy to a Biological Neuron:

- **Inputs:** Like dendrites receiving signals.
- **Weights:** Represent the strength or importance of each input signal (synaptic strength).
- **Summation:** The cell body (soma) integrates the incoming signals.
- **Activation Function (Threshold):** If the integrated signal exceeds a certain threshold, the neuron "fires" (sends an output signal down its axon).



3. Components of a Perceptron:

- **Inputs (x_1, x_2, \dots, x_n):** A vector of numerical features representing an instance.
- **Weights (w_1, w_2, \dots, w_n):** A vector of real-valued numbers, one for each input feature. These are learned during training.
- **Bias (b):** A special weight that is not associated with any input feature (or can be thought of as a weight for an input that is always +1). It allows shifting the decision boundary away from the origin, increasing model flexibility.
- **Summation Function (Net Input):** Calculates the weighted sum of the inputs plus the bias.
 - Mathematically: $z = (\sum_{i=1}^n w_i x_i) + b = \mathbf{w} \cdot \mathbf{x} + b$
- **Activation Function (Step Function):** Determines the output of the Perceptron based on the net input. The classical Perceptron uses a **Heaviside step function**:
 - Output is 1 if $z \geq \theta$ (threshold).
 - Output is 0 (or sometimes -1) if $z < \theta$.
 - Often, the threshold θ is incorporated into the bias term, so the activation function becomes:
 - Output is 1 if $z \geq 0$.
 - Output is 0 (or -1) if $z < 0$.

How a Perceptron Works (Forward Pass):

1. **Receive Inputs:** The Perceptron takes an input vector x .
2. **Compute Weighted Sum:** Each input x_i is multiplied by its corresponding weight w_i . These products are summed up, and the bias b is added: $z = w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b$
3. **Apply Activation Function:** The result z is passed through a step activation function to produce the final output y_{pred} (typically 0 or 1 for binary classification).
 - $y_{pred} = 1$ if $z \geq 0$
 - $y_{pred} = 0$ (or -1) if $z < 0$
4. This output y_{pred} is the Perceptron's prediction for the given input.

5. The Perceptron Learning Algorithm:

The goal of the learning algorithm is to find the optimal weights (w) and bias (b) that correctly classify the training examples.

1. **Initialization:**
 - Initialize weights (w) and bias (b) to small random values or zeros.
 - Define a **learning rate (η)**: A small positive value (e.g., 0.01, 0.1) that controls the step size of weight updates.
2. **Iterate through Training Data (Epochs):** For each training example $(x(j), d(j))$ where $x(j)$ is the input vector and $d(j)$ is the desired (true) output label (e.g., 0 or 1):
 - **a. Calculate Predicted Output:**
 - Compute the net input: $z(j) = \mathbf{w} \cdot \mathbf{x}(j) + b$
 - Apply the step function to get the predicted output $y(j)$.

- **b. Calculate Error:**
 - The error is the difference between the desired output and the predicted output: $error(j) = d(j) - y(j)$
- **c. Update Weights and Bias:** If $error(j) \neq 0$ (i.e., the prediction is wrong):
 - Weight update rule: For each weight w_i : $w_i(\text{new}) = w_i(\text{old}) + \eta \cdot error^{(j)} \cdot x_i^{(j)}$
 - Bias update rule: $b(\text{new}) = b(\text{old}) + \eta \cdot error^{(j)}$
 - If $error^{(j)} = 0$ (prediction is correct), no update is needed for that example.
- 3. **Repeat:** Repeat step 2 for a fixed number of epochs (passes through the entire training dataset) or until all training examples are correctly classified, or some other stopping criterion is met.

Intuition behind the update rule:

If $d^{(j)} = 1$ and $y^{(j)} = 0$ (false negative, error = 1): Weights are increased for positive inputs $x_i^{(j)}$, pushing z higher.

If $d^{(j)} = 0$ and $y^{(j)} = 1$ (false positive, error = -1): Weights are decreased for positive inputs $x_i^{(j)}$, pushing z lower.

- The learning rate η scales the magnitude of the update.

Convergence: The Perceptron learning algorithm is guaranteed to converge and find a solution (a separating hyperplane) **if and only if** the training data is **linearly separable**. If the data is not linearly separable, the algorithm will typically not converge and may oscillate.

6. Mathematics Summary:

- **Input Vector:** $\mathbf{x} = [x_1, x_2, \dots, x_n]$
- **Weight Vector:** $\mathbf{w} = [w_1, w_2, \dots, w_n]$
- **Bias:** b
- **Net Input (Weighted Sum):** $z = \mathbf{w}^T \mathbf{x} + b = \sum_{i=1}^n w_i x_i + b$
- **Activation (Output y for binary classification, labels 0 and 1)**
- **Error (for desired output d and predicted output y):** $e = d - y$
- **Weight Update:** $w_i \leftarrow w_i + \eta \cdot e \cdot x_i$
- **Bias Update:** $b \leftarrow b + \eta \cdot e$

Geometrically, the equation $\mathbf{w} \cdot \mathbf{x} + b = 0$ defines a hyperplane that separates the input space into two regions.


7. Limitations of the Perceptron:

- **Linear Separability:** The most significant limitation is that a single Perceptron can only classify data that is **linearly separable**. It cannot solve non-linearly separable problems, famously illustrated by the **XOR problem**.

- **Binary Output:** The classical Perceptron with a step function produces only binary outputs (0 or 1, or -1 or 1). It doesn't provide probability scores.
- **Sensitivity to Learning Rate:** If the learning rate is too high, the algorithm might overshoot the optimal weights and oscillate. If too low, convergence can be very slow.
- **Noisy Data:** Performance can degrade with noisy data or outliers.
- **Convergence on Non-Separable Data:** If the data is not linearly separable, the Perceptron learning rule will not converge; the weights will continue to change.

8. Significance & Legacy:

- **Foundation of Neural Networks:** Despite its limitations, the Perceptron was a crucial first step and forms the basic building block (neuron) for more complex, multi-layer neural networks (Multi-Layer Perceptrons or MLPs).



③ Perceptron [Artificial Neuron or Neural Network Unit]

① Input layer ✓ [Single Layered NN]

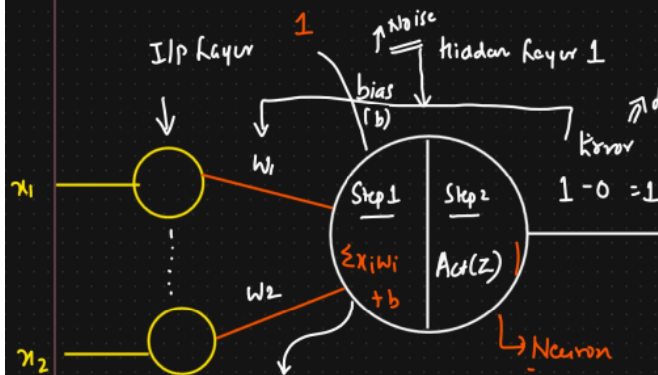
② Hidden layer ✓

③ Weights ✓

④ Activation function ✓

Binary classifier

DATASET



| x_1 IQ | x_2 No. of study hours | o/p Pass/Fail |
|-------------|-----------------------------|---------------|
| → 95 | 3 | 0 |
| → 110 | 4 | 1 |
| → 100 | 5 | 1 |

Activation function

Transform the output


between 0 to 1

-1 to 1

$$Z = w_1x_1 + w_2x_2 + b$$


$$Z = \sum_{i=1}^n x_i w_i + b$$

Step Function




$$\begin{cases} 0 & z \leq 0 \\ 1 & z > 0 \end{cases}$$

Sigmoid Function



$$\begin{cases} 1 & z > 0.5 \\ 0 & z \leq 0.5 \end{cases}$$



Step 1

$$z = \sum_{i=1}^n w_i x_i + b$$

$$z = b + w_1 x_1 + w_2 x_2 + w_3 x_3 + \dots + w_n x_n$$

↓

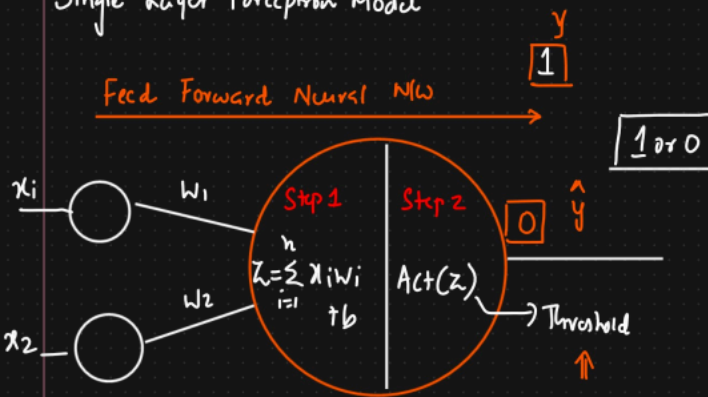
$$y = mx + c$$

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$

Linear
problem
statement

Perceptron Models

Single Layer Perceptron Model



Multi Layered Perceptron Model

- ① Forward Propagation
- ② Backward "
- ③ Loss functions
- ④ Activation functions
- ⑤ Optimizers

[ANN]

↓

linear separable problem



Binary classification



linearly separable



Non linear separable

⇒ Multi Layered NN